

# Data wrangling intro

## Contents

<b>1</b>	<b>Data</b>	<b>1</b>
1.1	Data example . . . . .	1
<b>2</b>	<b>Data wrangling</b>	<b>2</b>
2.1	The tidyverse package . . . . .	2
2.2	Selecting columns/variables . . . . .	2
2.3	Renaming columns/variables . . . . .	2
2.4	Filtering rows/cases/observations . . . . .	3
2.5	Arranging rows/cases/observations . . . . .	4
2.6	Extracting rows/cases/observations . . . . .	5
2.7	Extending (mutating) data with new variables . . . . .	6
2.8	Summarising data . . . . .	6
2.9	Summarising data by group . . . . .	6
2.10	Counting . . . . .	7
2.11	Bar graphs . . . . .	7
2.12	Exercises . . . . .	9

## 1 Data

### 1.1 Data example

We use data about penguins from the R package palmerpenguins

```
pengu <- palmerpenguins::penguins
pengu
```

```
## # A tibble: 344 x 8
##   species island   bill_length_mm bill_depth_mm flipper_length_mm body_mass_g
##   <fct>   <fct>         <dbl>         <dbl>         <int>         <int>
## 1 Adelie  Torgersen         39.1          18.7          181          3750
## 2 Adelie  Torgersen         39.5          17.4          186          3800
## 3 Adelie  Torgersen         40.3           18          195          3250
## 4 Adelie  Torgersen          NA           NA           NA           NA
## 5 Adelie  Torgersen         36.7          19.3          193          3450
## 6 Adelie  Torgersen         39.3          20.6          190          3650
## 7 Adelie  Torgersen         38.9          17.8          181          3625
## 8 Adelie  Torgersen         39.2          19.6          195          4675
## 9 Adelie  Torgersen         34.1          18.1          193          3475
## 10 Adelie Torgersen         42           20.2          190          4250
## # i 334 more rows
## # i 2 more variables: sex <fct>, year <int>
```

## 2 Data wrangling

### 2.1 The tidyverse package

We will use the tidyverse package extensively, and load it at the beginning of most sessions:

```
library(tidyverse)
```

```
## -- Attaching core tidyverse packages ----- tidyverse 2.0.0 --
## v dplyr      1.1.4      v readr      2.1.5
## v forcats    1.0.0      v stringr   1.5.1
## v ggplot2    3.5.1      v tibble    3.2.1
## v lubridate  1.9.3      v tidyr     1.3.1
## v purrr      1.0.2
## -- Conflicts ----- tidyverse_conflicts() --
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()     masks stats::lag()
## i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to become errors
```

### 2.2 Selecting columns/variables

- To subset columns of data use `select()` (automatically loaded from `dplyr` package by `tidyverse`):

```
bill_data <- select(pengu, bill_length_mm, bill_depth_mm,
                    species, sex, island, year)
```

- This particular subset can be written shorter by (overwriting the object we just created):

```
bill_data <- select(pengu, -flipper_length_mm, -body_mass_g)
```

- Special role of first argument gives rise to this “pipe” (`|>`) syntax:

```
bill_data <- pengu |> select(-flipper_length_mm, -body_mass_g)
```

- We read this as: first take the dataset **pengu** and **then** select all columns except `flipper_length_mm` and `body_mass_g`.
- The resulting dataset doesn’t have flipper length and body mass:

```
names(bill_data)
```

```
## [1] "species"      "island"        "bill_length_mm" "bill_depth_mm"
## [5] "sex"          "year"
```

### 2.3 Renaming columns/variables

Either via `select()` (which can also use column position), but then you only keep the selected ones:

```
bill_data_small <- bill_data |>
  select(len = bill_length_mm, depth = bill_depth_mm, pengu_type = 1)
bill_data_small
```

```
## # A tibble: 344 x 3
##   len depth pengu_type
##   <dbl> <dbl> <fct>
## 1  39.1  18.7 Adelie
## 2  39.5  17.4 Adelie
## 3  40.3   18  Adelie
## 4   NA   NA  Adelie
## 5  36.7  19.3 Adelie
```

```
## 6 39.3 20.6 Adelie
## 7 38.9 17.8 Adelie
## 8 39.2 19.6 Adelie
## 9 34.1 18.1 Adelie
## 10 42 20.2 Adelie
## # i 334 more rows
```

Or via `rename()` where you keep everything in place and just rename the relevant columns:

```
bill_data |> rename(len = bill_length_mm, depth = bill_depth_mm, pengu_type = 1)
```

```
## # A tibble: 344 x 6
##   pengu_type island      len depth sex      year
##   <fct>      <fct>    <dbl> <dbl> <fct> <int>
## 1 Adelie    Torgersen  39.1  18.7 male   2007
## 2 Adelie    Torgersen  39.5  17.4 female 2007
## 3 Adelie    Torgersen  40.3  18   female 2007
## 4 Adelie    Torgersen  NA    NA    <NA>   2007
## 5 Adelie    Torgersen  36.7  19.3 female 2007
## 6 Adelie    Torgersen  39.3  20.6 male   2007
## 7 Adelie    Torgersen  38.9  17.8 female 2007
## 8 Adelie    Torgersen  39.2  19.6 male   2007
## 9 Adelie    Torgersen  34.1  18.1 <NA>   2007
## 10 Adelie   Torgersen  42    20.2 <NA>   2007
## # i 334 more rows
```

## 2.4 Filtering rows/cases/observations

- We use `filter()` to subset rows/cases. E.g. all penguins from Biscoe islands:

```
pengu |> filter(island == "Biscoe")
```

```
## # A tibble: 168 x 8
##   species island bill_length_mm bill_depth_mm flipper_length_mm body_mass_g
##   <fct>    <fct>         <dbl>         <dbl>           <int>         <int>
## 1 Adelie  Biscoe           37.8           18.3             174           3400
## 2 Adelie  Biscoe           37.7           18.7             180           3600
## 3 Adelie  Biscoe           35.9           19.2             189           3800
## 4 Adelie  Biscoe           38.2           18.1             185           3950
## 5 Adelie  Biscoe           38.8           17.2             180           3800
## 6 Adelie  Biscoe           35.3           18.9             187           3800
## 7 Adelie  Biscoe           40.6           18.6             183           3550
## 8 Adelie  Biscoe           40.5           17.9             187           3200
## 9 Adelie  Biscoe           37.9           18.6             172           3150
## 10 Adelie Biscoe           40.5           18.9             180           3950
## # i 158 more rows
## # i 2 more variables: sex <fct>, year <int>
```

- All male Gentoo penguins with over 220 mm flippers:

```
pengu |> filter(sex == "male") |>
  filter(species == "Gentoo") |>
  filter(flipper_length_mm > 220)
```

```
## # A tibble: 34 x 8
##   species island bill_length_mm bill_depth_mm flipper_length_mm body_mass_g
```

```
##      <fct>   <fct>           <dbl>         <dbl>           <int>         <int>
## 1 Gentoo   Biscoe           50           16.3           230          5700
## 2 Gentoo   Biscoe          49.2           15.2           221          6300
## 3 Gentoo   Biscoe          48.7           15.1           222          5350
## 4 Gentoo   Biscoe          47.3           15.3           222          5250
## 5 Gentoo   Biscoe          59.6            17           230          6050
## 6 Gentoo   Biscoe          49.6            16           225          5700
## 7 Gentoo   Biscoe          50.5           15.9           222          5550
## 8 Gentoo   Biscoe          50.5           15.9           225          5400
## 9 Gentoo   Biscoe          50.1            15           225          5000
## 10 Gentoo  Biscoe          50.4           15.3           224          5550
## # i 24 more rows
## # i 2 more variables: sex <fct>, year <int>
```

- This could also have been done with a single `filter()` command (output not shown):

```
pengu |> filter(sex == "male" & species == "Gentoo" & flipper_length_mm > 220)
```

- 
- All penguins of species Gentoo or Adelie:

```
pengu |> filter(species == "Gentoo" | species == "Adelie")
```

```
## # A tibble: 276 x 8
##   species island bill_length_mm bill_depth_mm flipper_length_mm body_mass_g
##   <fct>   <fct>         <dbl>         <dbl>           <int>         <int>
## 1 Adelie  Torgersen         39.1          18.7           181          3750
## 2 Adelie  Torgersen         39.5          17.4           186          3800
## 3 Adelie  Torgersen         40.3           18           195          3250
## 4 Adelie  Torgersen          NA           NA            NA            NA
## 5 Adelie  Torgersen         36.7          19.3           193          3450
## 6 Adelie  Torgersen         39.3          20.6           190          3650
## 7 Adelie  Torgersen         38.9          17.8           181          3625
## 8 Adelie  Torgersen         39.2          19.6           195          4675
## 9 Adelie  Torgersen         34.1          18.1           193          3475
## 10 Adelie Torgersen         42           20.2           190          4250
## # i 266 more rows
## # i 2 more variables: sex <fct>, year <int>
```

- This would be the same as penguins which are not Chinstrap (output not shown):

```
pengu |> filter(species != "Chinstrap")
```

## 2.5 Arranging rows/cases/observations

- We use `arrange()` to arrange the order of the rows/cases:

```
pengu |> filter(sex == "female") |> arrange(body_mass_g)
```

```
## # A tibble: 165 x 8
##   species island bill_length_mm bill_depth_mm flipper_length_mm body_mass_g
##   <fct>   <fct>         <dbl>         <dbl>           <int>         <int>
## 1 Chinstrap Dream         46.9          16.6           192          2700
## 2 Adelie   Biscoe         36.5          16.6           181          2850
## 3 Adelie   Biscoe         36.4          17.1           184          2850
## 4 Adelie   Biscoe         34.5          18.1           187          2900
## 5 Adelie   Dream         33.1          16.1           178          2900
```

```
## 6 Adelie Torgers~ 38.6 17 188 2900
## 7 Chinstrap Dream 43.2 16.6 187 2900
## 8 Adelie Biscoe 37.9 18.6 193 2925
## 9 Adelie Dream 37 16.9 185 3000
## 10 Adelie Dream 37.3 16.8 192 3000
## # i 155 more rows
## # i 2 more variables: sex <fct>, year <int>
```

- Use `arrange(desc())` for descending values:

```
pengu |> filter(sex == "female") |> arrange(desc(body_mass_g))
```

```
## # A tibble: 165 x 8
##   species island bill_length_mm bill_depth_mm flipper_length_mm body_mass_g
##   <fct>   <fct>         <dbl>         <dbl>         <int>         <int>
## 1 Gentoo Biscoe         46.5          14.8          217          5200
## 2 Gentoo Biscoe         45.2          14.8          212          5200
## 3 Gentoo Biscoe         49.1          14.8          220          5150
## 4 Gentoo Biscoe         44.9          13.3          213          5100
## 5 Gentoo Biscoe         45.1          14.5          207          5050
## 6 Gentoo Biscoe         45.1          14.5          215          5000
## 7 Gentoo Biscoe         42.9          13.1          215          5000
## 8 Gentoo Biscoe         50.5          15.2          216          5000
## 9 Gentoo Biscoe         47.2          15.5          215          4975
## 10 Gentoo Biscoe         42.6          13.7          213          4950
## # i 155 more rows
## # i 2 more variables: sex <fct>, year <int>
```

## 2.6 Extracting rows/cases/observations

- The `arrange()` command reorders the rows, so the ones we are interested in are in the top or bottom. If we want to extract e.g. the top 5 heaviest female penguins we use `slice_max()`:

```
pengu |> filter(sex == "female") |> slice_max(body_mass_g, n = 5)
```

```
## # A tibble: 5 x 8
##   species island bill_length_mm bill_depth_mm flipper_length_mm body_mass_g
##   <fct>   <fct>         <dbl>         <dbl>         <int>         <int>
## 1 Gentoo Biscoe         46.5          14.8          217          5200
## 2 Gentoo Biscoe         45.2          14.8          212          5200
## 3 Gentoo Biscoe         49.1          14.8          220          5150
## 4 Gentoo Biscoe         44.9          13.3          213          5100
## 5 Gentoo Biscoe         45.1          14.5          207          5050
## # i 2 more variables: sex <fct>, year <int>
```

- To extract rows by their rownumber we simply use `slice()` (note the five female penguins below are taken from the original order of the dataset and not according to some ordering variable):

```
pengu |> filter(sex == "female") |> slice(c(1, 3, 5, 7, 9))
```

```
## # A tibble: 5 x 8
##   species island bill_length_mm bill_depth_mm flipper_length_mm body_mass_g
##   <fct>   <fct>         <dbl>         <dbl>         <int>         <int>
## 1 Adelie Torgersen     39.5          17.4          186          3800
## 2 Adelie Torgersen     36.7          19.3          193          3450
## 3 Adelie Torgersen     41.1          17.6          182          3200
## 4 Adelie Torgersen     38.7          19          195          3450
```

```
## 5 Adelie Biscoe          37.8          18.3          174          3400
## # i 2 more variables: sex <fct>, year <int>
```

- Other slicing command are: `slice_min()`, `slice_sample()`, `slice_head()` and `slice_tail()`.

## 2.7 Extending (mutating) data with new variables

```
bill_data_ratio <- bill_data_small |> mutate(ratio = len/depth)
bill_data_ratio
```

```
## # A tibble: 344 x 4
##   len depth pengu_type ratio
##   <dbl> <dbl> <fct>    <dbl>
## 1  39.1  18.7 Adelie      2.09
## 2  39.5  17.4 Adelie      2.27
## 3  40.3  18   Adelie      2.24
## 4  NA    NA   Adelie      NA
## 5  36.7  19.3 Adelie      1.90
## 6  39.3  20.6 Adelie      1.91
## 7  38.9  17.8 Adelie      2.19
## 8  39.2  19.6 Adelie      2
## 9  34.1  18.1 Adelie      1.88
## 10 42    20.2 Adelie      2.08
## # i 334 more rows
```

## 2.8 Summarising data

```
bill_data_small |> summarise(avg_len = mean(len), avg_depth = mean(depth),
                             sd_len = sd(len), sd_depth = sd(depth),
                             cor_len_depth = cor(len, depth))
```

```
## # A tibble: 1 x 5
##   avg_len avg_depth sd_len sd_depth cor_len_depth
##   <dbl>    <dbl> <dbl>    <dbl>    <dbl>
## 1    NA        NA    NA        NA        NA
```

By default R cannot calculate summaries of data with NA, so we can either tell each function to remove NA or drop all NAs from the data (we do the latter):

```
bill_data_small |> drop_na() |>
  summarise(avg_len = mean(len), avg_depth = mean(depth),
            sd_len = sd(len), sd_depth = sd(depth),
            cor_len_depth = cor(len, depth))
```

```
## # A tibble: 1 x 5
##   avg_len avg_depth sd_len sd_depth cor_len_depth
##   <dbl>    <dbl> <dbl>    <dbl>    <dbl>
## 1  43.9      17.2  5.46     1.97     -0.235
```

## 2.9 Summarising data by group

The `group_by()` function simply annotates the data as being grouped:

```
bill_grp <- bill_data_small |> group_by(pengu_type)
bill_grp
```

```
## # A tibble: 344 x 3
## # Groups:   pengu_type [3]
##   len depth pengu_type
##   <dbl> <dbl> <fct>
## 1  39.1  18.7 Adelie
## 2  39.5  17.4 Adelie
## 3  40.3   18  Adelie
## 4   NA    NA  Adelie
## 5  36.7  19.3 Adelie
## 6  39.3  20.6 Adelie
## 7  38.9  17.8 Adelie
## 8  39.2  19.6 Adelie
## 9  34.1  18.1 Adelie
## 10 42    20.2 Adelie
## # i 334 more rows
```

Then summaries are done for each group:

```
bill_data_small |> na.omit() |> group_by(pengu_type) |>
  summarise(avg_len = mean(len), avg_depth = mean(depth),
            sd_len = sd(len), sd_depth = sd(depth),
            cor_len_depth = cor(len, depth))
```

```
## # A tibble: 3 x 6
##   pengu_type avg_len avg_depth sd_len sd_depth cor_len_depth
##   <fct>      <dbl>    <dbl> <dbl>    <dbl>        <dbl>
## 1 Adelie      38.8      18.3   2.66    1.22         0.391
## 2 Chinstrap   48.8      18.4   3.34    1.14         0.654
## 3 Gentoo     47.5      15.0   3.08    0.981         0.643
```

## 2.10 Counting

Counting is a very common summary of data and it has a separate function/verb `count` which is effectively just shorthand for `group_by()` + `summarise(n = n())`:

```
pengu |> count(species)
```

```
## # A tibble: 3 x 2
##   species      n
##   <fct>    <int>
## 1 Adelie   152
## 2 Chinstrap 68
## 3 Gentoo  124
```

```
pengu |> group_by(species) |> summarise(n = n())
```

```
## # A tibble: 3 x 2
##   species      n
##   <fct>    <int>
## 1 Adelie   152
## 2 Chinstrap 68
## 3 Gentoo  124
```

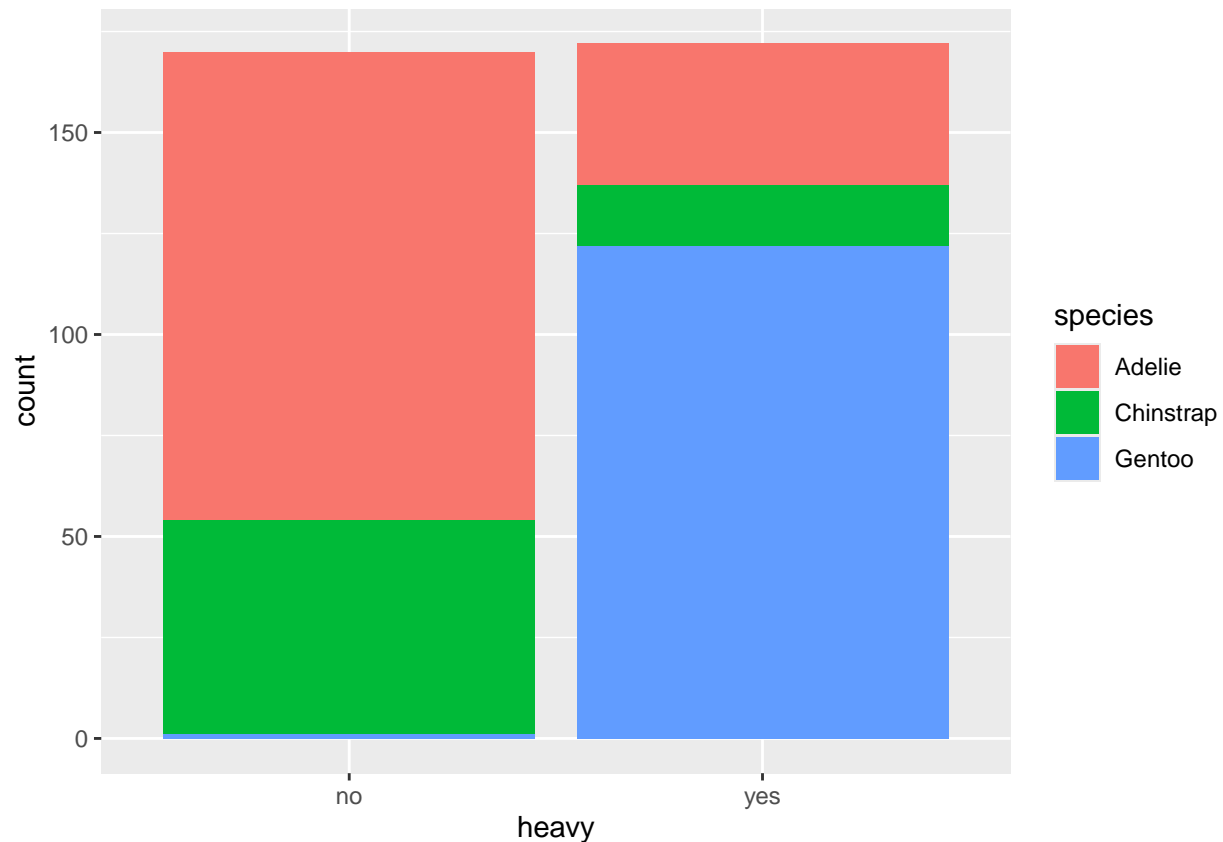
## 2.11 Bar graphs

To make a bar graph from the “raw” data where we didn’t count things first use `geom_bar()`. Before making the graph we make a new variable `heavy` indicating whether the penguin mass is above 4000 g (the command

`ifelse()` checks the condition given as first input argument and if the condition is `TRUE` it returns the one option and if it is false it returns the other (these options are given as second and third input arguments, respectively):

```
pengu3 <- pengu |>
  drop_na(body_mass_g) |>
  mutate(heavy = ifelse(body_mass_g > 4000, "yes", "no"))
```

```
ggplot(pengu3, aes(x = heavy, fill = species)) +
  geom_bar()
```



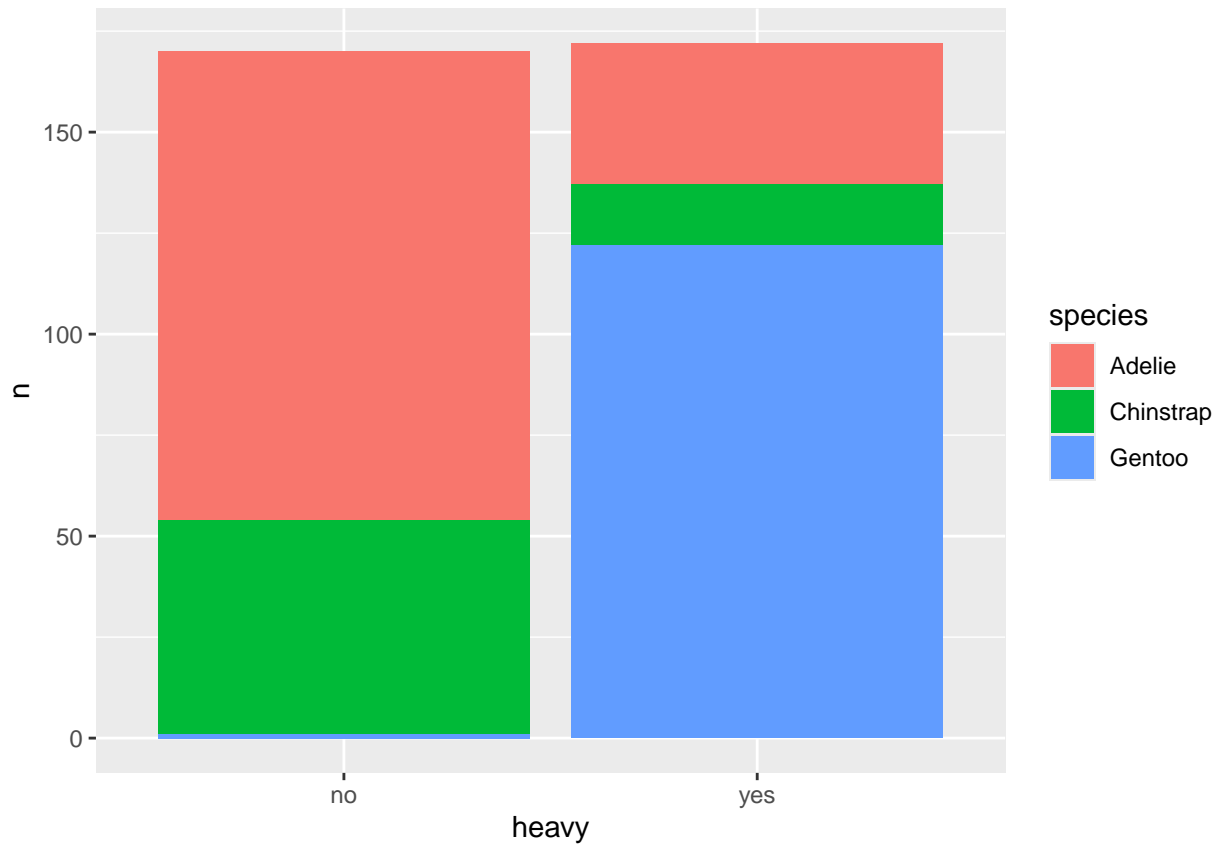
To make a bar graph from the aggregated data where we have the counts we want to plot use `geom_col()` (for column):

```
pengu3count <- pengu3 |>
  count(heavy, species)
pengu3count
```

```
## # A tibble: 6 x 3
##   heavy species      n
##   <chr> <fct>    <int>
## 1 no    Adelie     116
## 2 no    Chinstrap   53
## 3 no    Gentoo      1
## 4 yes   Adelie      35
## 5 yes   Chinstrap   15
## 6 yes   Gentoo     122
```



```
ggplot(pengu3count, aes(x = heavy, y = n, fill = species)) +  
  geom_col()
```



## 2.12 Exercises

- First work on the qmd exercise on Moodle.
- If you have extra time confront Chapter 3 of R4DS and work on (you will need to read about the NYC flights data to understand the exercises):
  - Exercises 3.2.5: <https://r4ds.hadley.nz/data-transform#exercises>
  - Exercises 3.3.5: <https://r4ds.hadley.nz/data-transform#exercises-1>
  - Exercises 3.5.7: <https://r4ds.hadley.nz/data-transform#exercises-2>