

# Classification and tree-based methods

## Introduction

### Terminology

- We will consider Classification and regression trees (CART) as well as the related Random Forests
  - These are methods of supervised learning (*“labelled”* training data) for:
    - Classification
    - Regression
  - We focus on classification
- 

### Classification

- Given a feature vector  $x$  and a qualitative response  $Y$  taking values in the set  $C$ , the classification task is to build a function  $f(x)$  that takes as input the feature vector  $x$  and predicts its value for  $Y$ ; i.e.  $f(x) \in C$
- Often: interested in estimating the probabilities that  $X$  belongs to each category in  $C$

Many methods for classification:

- Logistic regression
  - Classification (and regression) trees
  - Random Forest
  - Nearest Neighbours
  - Naive Bayes
  - Support Vector Machines (SVM)
  - Neural Networks
  - ...
- 

## Classification trees

One limitation of logistic regression (and multinomial regression) is the need for a statistical model, with e.g. the assumption logit-linear relationships, which typically is the biggest challenge - the relationship between the probability of the outcome and the explanatory variables may be anything but linear on the (arbitrary) logit-scale.

A very generic, assumption free, adaptable and flexible class of models are the classification and regression trees. The seminal CART-book (Classification And Regression Trees) from 1984 by Breiman et al. laid much of the foundation for their success.

---

## Partition concept with small example

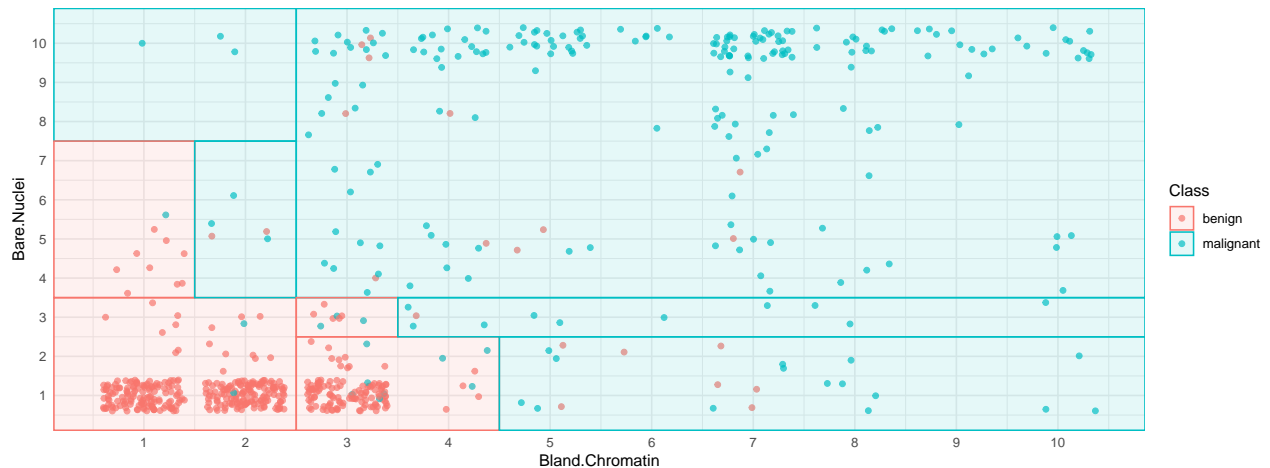
```
#' Plot uses non-standard package installable from r-universe.dev:
#' install.packages('parttree', repos = 'https://grantmcdermott.r-universe.dev')
BCfull <- na.omit(OneR::breastcancer) |>
  as_tibble(.name_repair = "universal")
```

New names:

```
* `Clump Thickness` -> `Clump.Thickness`
* `Uniformity of Cell Size` -> `Uniformity.of.Cell.Size`
* `Uniformity of Cell Shape` -> `Uniformity.of.Cell.Shape`
* `Marginal Adhesion` -> `Marginal.Adhesion`
* `Single Epithelial Cell Size` -> `Single.Epithelial.Cell.Size`
* `Bare Nuclei` -> `Bare.Nuclei`
* `Bland Chromatin` -> `Bland.Chromatin`
* `Normal Nucleoli` -> `Normal.Nucleoli`
```

```
BC_first_tree <- rpart::rpart(Class ~ Bland.Chromatin + Bare.Nuclei, data = BCfull, cp = 0)
BCfull |>
```

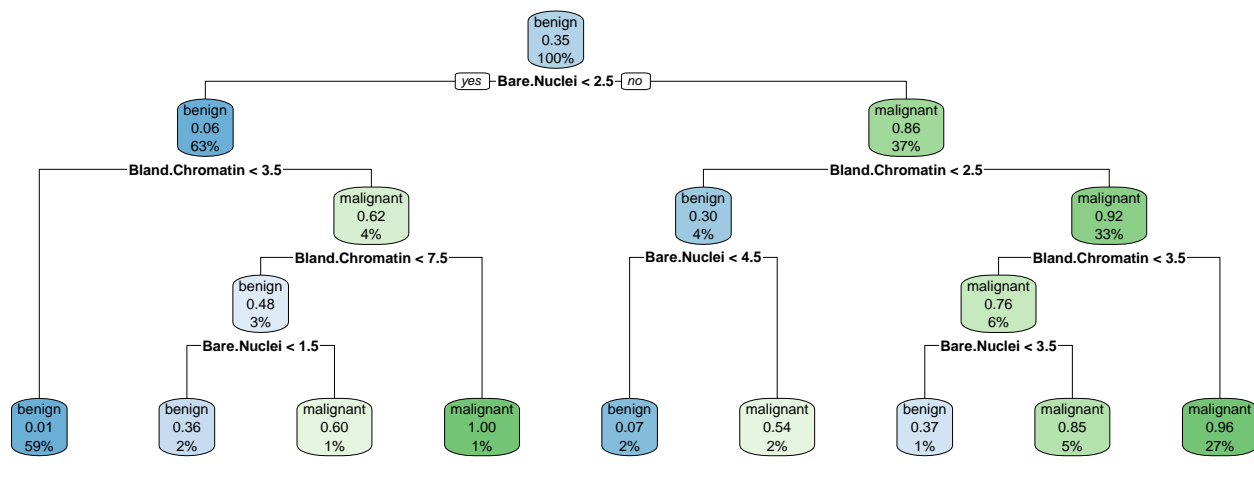
```
  ggplot(aes(x = Bland.Chromatin, y = Bare.Nuclei, color = Class)) +
  geom_jitter(alpha=0.7) +
  parttree::geom_parttree(data = BC_first_tree, aes(fill=Class), alpha = 0.1) +
  theme_minimal() +
  scale_y_continuous(n.breaks = 10) +
  scale_x_continuous(n.breaks = 10)
```



```
library(rpart.plot)
```

Loading required package: rpart

```
rpart.plot(BC_first_tree, roundint = FALSE)
```



## rpart

For plotting the `rpart.plot` package is excellent: <http://www.milbo.org/rpart-plot/prp.pdf>

For illustration purposes we set parameters to get a big tree which is clearly over-fitting the data

```
library(rpart)
set.seed(42)
BC_rpart <- rpart(Class ~ ., data = BCfull, cp = 0, minsplit = 2, minbucket = 1)
```

```
BC_rpart
```

```
n= 683
```

```
node), split, n, loss, yval, (yprob)
* denotes terminal node
```

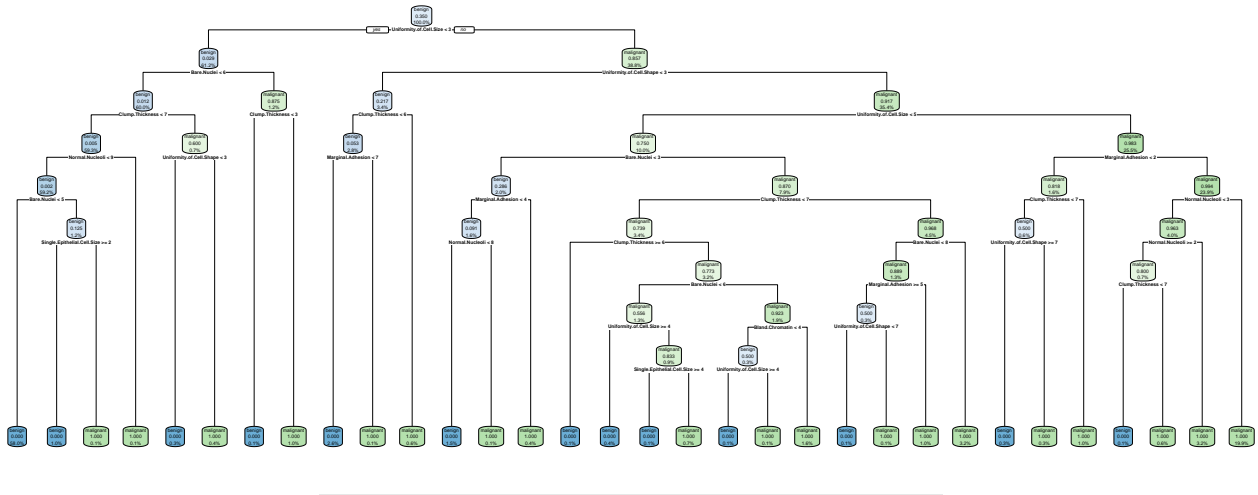
```
1) root 683 239 benign (0.650073206 0.349926794)
  2) Uniformity.of.Cell.Size< 2.5 418 12 benign (0.971291866 0.028708134)
    4) Bare.Nuclei< 5.5 410 5 benign (0.987804878 0.012195122)
      8) Clump.Thickness< 6.5 405 2 benign (0.995061728 0.004938272)
        16) Normal.Nucleoli< 9 404 1 benign (0.997524752 0.002475248)
          32) Bare.Nuclei< 4.5 396 0 benign (1.000000000 0.000000000) *
            33) Bare.Nuclei>=4.5 8 1 benign (0.875000000 0.125000000)
              66) Single.Epithelial.Cell.Size>=1.5 7 0 benign (1.000000000 0.000000000) *
                67) Single.Epithelial.Cell.Size< 1.5 1 0 malignant (0.000000000 1.000000000) *
              17) Normal.Nucleoli>=9 1 0 malignant (0.000000000 1.000000000) *
            9) Clump.Thickness>=6.5 5 2 malignant (0.400000000 0.600000000)
              18) Uniformity.of.Cell.Shape< 2.5 2 0 benign (1.000000000 0.000000000) *
                19) Uniformity.of.Cell.Shape>=2.5 3 0 malignant (0.000000000 1.000000000) *
            5) Bare.Nuclei>=5.5 8 1 malignant (0.125000000 0.875000000)
              10) Clump.Thickness< 2.5 1 0 benign (1.000000000 0.000000000) *
                11) Clump.Thickness>=2.5 7 0 malignant (0.000000000 1.000000000) *
          3) Uniformity.of.Cell.Size>=2.5 265 38 malignant (0.143396226 0.856603774)
            6) Uniformity.of.Cell.Shape< 2.5 23 5 benign (0.782608696 0.217391304)
              12) Clump.Thickness< 5.5 19 1 benign (0.947368421 0.052631579)
                24) Marginal.Adhesion< 7 18 0 benign (1.000000000 0.000000000) *
                  25) Marginal.Adhesion>=7 1 0 malignant (0.000000000 1.000000000) *
                13) Clump.Thickness>=5.5 4 0 malignant (0.000000000 1.000000000) *
```

```

7) Uniformity.of.Cell.Shape>=2.5 242 20 malignant (0.082644628 0.917355372)
14) Uniformity.of.Cell.Size< 4.5 68 17 malignant (0.250000000 0.750000000)
28) Bare.Nuclei< 2.5 14 4 benign (0.714285714 0.285714286)
56) Marginal.Adhesion< 3.5 11 1 benign (0.909090909 0.090909091)
112) Normal.Nucleoli< 7.5 10 0 benign (1.000000000 0.000000000) *
113) Normal.Nucleoli>=7.5 1 0 malignant (0.000000000 1.000000000) *
57) Marginal.Adhesion>=3.5 3 0 malignant (0.000000000 1.000000000) *
29) Bare.Nuclei>=2.5 54 7 malignant (0.129629630 0.870370370)
58) Clump.Thickness< 6.5 23 6 malignant (0.260869565 0.739130435)
116) Clump.Thickness>=5.5 1 0 benign (1.000000000 0.000000000) *
117) Clump.Thickness< 5.5 22 5 malignant (0.227272727 0.772727273)
234) Bare.Nuclei< 6 9 4 malignant (0.444444444 0.555555556)
468) Uniformity.of.Cell.Size>=3.5 3 0 benign (1.000000000 0.000000000) *
469) Uniformity.of.Cell.Size< 3.5 6 1 malignant (0.166666667 0.833333333)
938) Single.Epithelial.Cell.Size>=3.5 1 0 benign (1.000000000 0.000000000) *
939) Single.Epithelial.Cell.Size< 3.5 5 0 malignant (0.000000000 1.000000000) *
235) Bare.Nuclei>=6 13 1 malignant (0.076923077 0.923076923)
470) Bland.Chromatin< 3.5 2 1 benign (0.500000000 0.500000000)
940) Uniformity.of.Cell.Size>=3.5 1 0 benign (1.000000000 0.000000000) *
941) Uniformity.of.Cell.Size< 3.5 1 0 malignant (0.000000000 1.000000000) *
471) Bland.Chromatin>=3.5 11 0 malignant (0.000000000 1.000000000) *
59) Clump.Thickness>=6.5 31 1 malignant (0.032258065 0.967741935)
118) Bare.Nuclei< 7.5 9 1 malignant (0.111111111 0.888888889)
236) Marginal.Adhesion>=4.5 2 1 benign (0.500000000 0.500000000)
472) Uniformity.of.Cell.Shape< 7 1 0 benign (1.000000000 0.000000000) *
473) Uniformity.of.Cell.Shape>=7 1 0 malignant (0.000000000 1.000000000) *
237) Marginal.Adhesion< 4.5 7 0 malignant (0.000000000 1.000000000) *
119) Bare.Nuclei>=7.5 22 0 malignant (0.000000000 1.000000000) *
15) Uniformity.of.Cell.Size>=4.5 174 3 malignant (0.017241379 0.982758621)
30) Marginal.Adhesion< 1.5 11 2 malignant (0.181818182 0.818181818)
60) Clump.Thickness< 7 4 2 benign (0.500000000 0.500000000)
120) Uniformity.of.Cell.Shape>=6.5 2 0 benign (1.000000000 0.000000000) *
121) Uniformity.of.Cell.Shape< 6.5 2 0 malignant (0.000000000 1.000000000) *
61) Clump.Thickness>=7 7 0 malignant (0.000000000 1.000000000) *
31) Marginal.Adhesion>=1.5 163 1 malignant (0.006134969 0.993865031)
62) Normal.Nucleoli< 2.5 27 1 malignant (0.037037037 0.962962963)
124) Normal.Nucleoli>=1.5 5 1 malignant (0.200000000 0.800000000)
248) Clump.Thickness< 6.5 1 0 benign (1.000000000 0.000000000) *
249) Clump.Thickness>=6.5 4 0 malignant (0.000000000 1.000000000) *
125) Normal.Nucleoli< 1.5 22 0 malignant (0.000000000 1.000000000) *
63) Normal.Nucleoli>=2.5 136 0 malignant (0.000000000 1.000000000) *

```

```
rpart.plot(BC_rpart, digits = 3)
```



## Complexity?

```
printcp(BC_rpart)
```

Classification tree:

```
rpart(formula = Class ~ ., data = BCfull, cp = 0, minsplit = 2,
      minbucket = 1)
```

Variables actually used in tree construction:


[1] Bare.Nuclei	Bland.Chromatin
[3] Clump.Thickness	Marginal.Adhesion
[5] Normal.Nucleoli	Single.Epithelial.Cell.Size
[7] Uniformity.of.Cell.Shape	Uniformity.of.Cell.Size

Root node error: 239/683 = 0.34993

n= 683



	CP	nsplit	rel error	xerror	xstd
1	0.7907950	0	1.0000000	1.00000	0.052153
2	0.0543933	1	0.2092050	0.29289	0.033164
3	0.0251046	2	0.1548117	0.18410	0.026845
4	0.0167364	3	0.1297071	0.18410	0.026845
5	0.0125523	4	0.1129707	0.17992	0.026559
6	0.0062762	7	0.0753138	0.11715	0.021682
7	0.0041841	9	0.0627615	0.12552	0.022408
8	0.0027894	18	0.0251046	0.12134	0.022049
9	0.0020921	21	0.0167364	0.11715	0.021682
10	0.0013947	25	0.0083682	0.13808	0.023448
11	0.0000000	31	0.0000000	0.14644	0.024111

```
plotcp(BC_rpart)
```



```
cp_value <- 0.008 ## from printcp/plotcp output
BC_rpart_prune <- prune(BC_rpart, cp = cp_value)

par(mfrow = c(2,1))
rpart.plot(BC_rpart, main = "Full tree")
rpart.plot(BC_rpart_prune, main = paste("Pruned tree, cp:", cp_value))
```



```
par(mfrow = c(1,1))
```

## Prediction

```
BC_pred <- as_tibble(predict(BC_rpart_prune, type = "prob")) |>
  mutate(pred_class = ifelse(benign > .5, "benign", "malignant"))

BCfull |> select(Class) |> bind_cols(BC_pred) |>
  count(Class, pred_class)
```

```
# A tibble: 4 x 3
  Class pred_class    n
  <fct>   <chr>    <int>
1 benign  benign    433
2 benign  malignant  11
```

3 malignant benign	7
4 malignant malignant	232

---

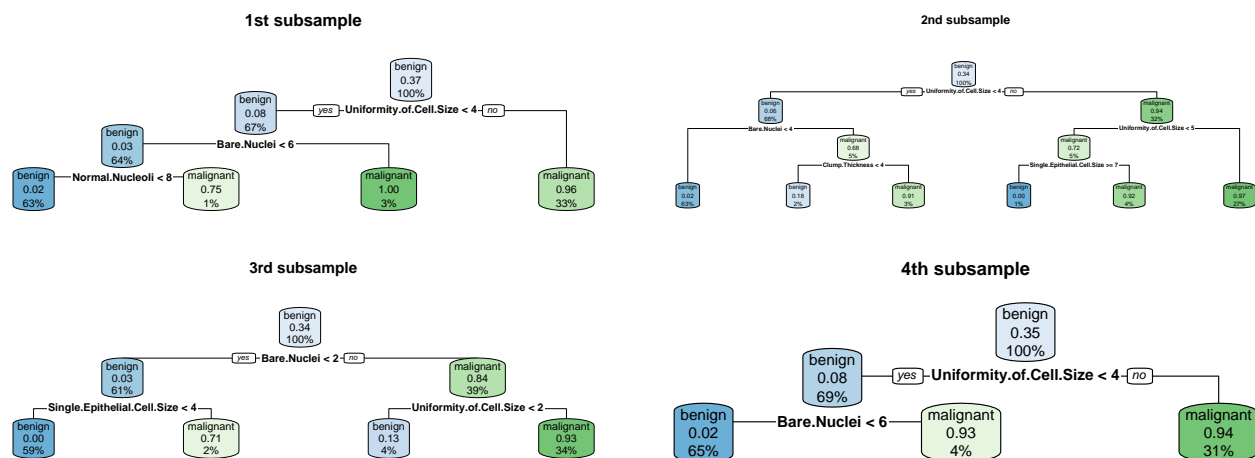
## Flexibility

Trees can be very non-robust. In other words, a small change in the data can cause a large change in the final estimated tree

```
set.seed(123)
N <- nrow(BCfull)
n_rep <- 4
row_index <- replicate(n_rep,
                        sample(N, size = N, replace = TRUE),
                        simplify = FALSE)

BC_resamp <-
  lapply(row_index, function(x) rpart(Class ~ ., data = BCfull, subset = x))

par(mfrow = c(2,2))
rpart.plot(BC_resamp[[1]], main = "1st subsample")
rpart.plot(BC_resamp[[2]], main = "2nd subsample")
rpart.plot(BC_resamp[[3]], main = "3rd subsample")
rpart.plot(BC_resamp[[4]], main = "4th subsample")
```



```
par(mfrow = c(1,1))
```

---

## Random forests

Random forest tries to remedy this flexibility of tree-based models by an ensemble of trees. The trees created on the previous slide are not identical, but still correlated in the sense that they use the same features/explanatory variables for splits and creating the trees. Furthermore, their predictions are highly correlated with each other.

```
set.seed(123)
BC_resamp |> lapply( function(x) predict(x, newdata = BCfull, type = "class")) |> set_names(paste0("res", 1:n_rep))
bind_cols(Class = BCfull$Class) |>
mutate(row = row_number()) |>
```

```
relocate(row, Class) |>
slice_sample(n = 10)
```

```
# A tibble: 10 x 6
  row Class    resample1 resample2 resample3 resample4
  <int> <fct>      <fct>      <fct>      <fct>      <fct>
1   415 benign    benign     benign     benign     benign
2   463 benign    benign     benign     benign     benign
3   179 malignant malignant malignant malignant malignant
4   526 benign    benign     benign     benign     benign
5   195 malignant malignant malignant malignant malignant
6   118 benign    benign     benign     malignant    benign
7   299 benign    benign     benign     benign     benign
8   229 benign    benign     benign     benign     benign
9   244 malignant malignant malignant malignant malignant
10   14 benign    benign     benign     benign     benign
```

---

## Uncorrelated trees

The success of random forest relies on the simple fact that for uncorrelated quantities, the average is a consistent and unbiased estimator with a variance going to zero as  $1/\#trees$ .

Because of the bootstrapping of the data, the non-included observations (called out-of-bag samples) are used in random forests to assess the accuracy of the model.

---

## Package randomForest

```
library(randomForest)
```

```
randomForest 4.7-1.1
```

Type `rfNews()` to see new features/changes/bug fixes.

Attaching package: 'randomForest'

The following object is masked from 'package:dplyr':

```
combine
```

The following object is masked from 'package:ggplot2':

```
margin
```

```
set.seed(1234)
```

```
BC_RF <- randomForest(Class ~ ., data = BCfull, importance=TRUE)
BC_RF
```

Call:

```
randomForest(formula = Class ~ ., data = BCfull, importance = TRUE)
  Type of random forest: classification
    Number of trees: 500
```



No. of variables tried at each split: 3

OOB estimate of error rate: 2.78%


Confusion matrix:

	benign	malignant	class.error
benign	432	12	0.02702703
malignant	7	232	0.02928870

## Variable importance for RF

`varImpPlot(BC_RF)`

BC\_RF




## Tuning of algorithm/hyper parameters of RF

Tuning on mtry with ntree fixed:

```
BC_RF_tune <- tuneRF(y = BCfull$Class, x = select(BCfull, -Class), improve = 0.001)
```

```
mtry = 3  OOB error = 3.07%
Searching left ...
mtry = 2   OOB error = 2.49%
0.1904762 0.001
mtry = 1   OOB error = 3.07%
-0.2352941 0.001
Searching right ...
mtry = 6   OOB error = 3.37%
-0.3529412 0.001
```




Tuning both mtry and ntree

```
tune_one <- function(n){
  output <- tuneRF(y = BCfull$Class, x = select(BCfull, -Class),
                    ntreeTry = n, improve = 0.05, trace = FALSE, plot = FALSE)
  return(bind_cols(nTree = n, output))
}
nTree <- c(10, 50, 100, 200, 500, 1000)
OOB_error <- map(nTree, tune_one)
```

```
-0.05951958 0.05
-0.1295392 0.05
-0.1052632 0.05
-0.3157895 0.05
-0.04545455 0.05
-0.1363636 0.05
0.1428571 0.05
-0.1111111 0.05
-0.2777778 0.05
-0.1111111 0.05
-0.3333333 0.05
0 0.05
-0.2777778 0.05
```

```
OOB_error |>
  bind_rows() |>
  ggplot(aes(x = mtry, y = OOBError, colour = factor(nTree))) +
  geom_point() + geom_line()
```



## Test and training set for RF

We can also use the `test` arguments of `randomForest` in order to assess the accuracy on a test set while fitting the model.

```
train_id <- sample(nrow(BCfull), 600)

BC_train <- BCfull[train_id,]
BC_test  <- BCfull[-train_id,]

BC_RF <- randomForest(Class ~ ., data = BC_train, importance=TRUE,
                      xtest = select(BC_test, - Class), ytest = BC_test$Class)

RF_test_error <- function(rf){
  rf_conf <- rf$test$confusion
  1 - sum(diag(rf_conf))/sum(rf_conf[, -ncol(rf_conf)])
}

p <- ncol(BC_train)-1

BC_RF_test <-
  expand_grid(
    nTree = c(1, seq(from = 20, to = 800, by = 20)),
    mTry = c(p, p/2, sqrt(p)) |>
    mutate(
      test_error = map2_dbl(.x = nTree, .y = mTry, \(n,m) RF_test_error(
        randomForest(Class ~ ., data = BC_train, importance=FALSE, ntree = n, mtry = m,
                      xtest = select(BC_test, - Class), ytest = BC_test$Class)
      ))
    )

BC_RF_test |>
  mutate(
    m = case_when(mTry == p ~ "p",
                  mTry == p/2 ~ "p/2",
                  TRUE ~ "sqrt(p)",
    ) |>
```

```
ggplot(aes(nTree, test_error, colour = m)) +  
  geom_line()
```

